



Programming meets Mathematics: **Optimization**

Sunglok Choi, Assistant Professor, Ph.D.
Computer Science and Engineering Department, SeoulTech
sunglok@seoultech.ac.kr | <https://mint-lab.github.io/>

Programming meets Mathematics

~~■ Calculus~~ **Differentiation**

~~■ Linear Algebra~~ **Vector and Matrix**

- [NumPy](#): `numpy.array` vs. `list/tuple`
- Vector: Why?
 - Vector multiplication: Dot product, cross product
- Matrix: Why?
 - Matrix multiplication
 - Matrix inverse (square + full rank), pseudo-inverse
 - Examples) Line and curve fitting (solving a system of linear equations)

■ **Optimization**

■ **Probability**

■ **Information Theory**

Getting Started from Line Fitting

(1, 4)

(x_i, y_i)

(4, 2)

- Line representation: $y = wx + b$ ($y = -\frac{2}{3}x + \frac{14}{3}$)
- Slope $w = \frac{2-4}{4-1} = -\frac{2}{3}$
- Y intercept $b = 4 - m \cdot 1 = \frac{14}{3}$
- Algebraic distance $d_a = y - (wx_i + b)$
(signed distance)

How to measure distance between a point and a line?

Getting Started from Line Fitting

(1, 4)



(x_i, y_i)



(4, 2)



- Line representation: $ax + by + c = 0$
($2x + 3y - 14 = 0$; $4x + 6y - 28 = 0$)
- Its shorter form: $\mathbf{n}^T \mathbf{x} + c = 0$
($\mathbf{n} = \begin{bmatrix} a \\ b \end{bmatrix}$ and $\mathbf{x} = \begin{bmatrix} x \\ y \end{bmatrix}$)
- Geometric distance $d_g = \frac{ax_i + by_i + c}{\sqrt{a^2 + b^2}} = \frac{\mathbf{n}^T \mathbf{x}_i + c}{\|\mathbf{n}\|}$
(signed distance)


Normal vector



How to measure distance between a **point** and a **line**?


Optimization

- **Optimization** is the selection of best element, with regard to some **criterion**, from a defined domain.
 - Alias: Mathematical programming
 - [Linear programming](#), [convex programming](#), [nonlinear programming](#), ..., [dynamic programming](#)
 - In the simplest case, optimization is maximizing or minimizing a **objective function**
 - Maximization: Objective functions \rightarrow profit/utility/fitness/reward/... functions
 - Minimization: Objective functions \rightarrow loss/cost/error/penalty/... functions
 - cf. Maximization and minimization is dual. \rightarrow Minimization is usually preferred.
 - Example) Finding x and y for the maximum z with $z = 4 - (x^2 + y^2)$
 - Unknown variable: $\mathbf{x} = [x, y]$ and its domain \mathbb{R}^2
 - Objective function: $f(x, y) = 4 - (x^2 + y^2)$ as a maximization problem
 - In short, $\hat{\mathbf{x}} = \underset{\mathbf{x}}{\operatorname{argmax}} f(\mathbf{x})$ where $\mathbf{x} \in \mathbb{R}^2$ and $f(\mathbf{x}) = 4 - \|\mathbf{x}\|_2^2$
 - cf. 2-[norm](#) (Euclidean [norm](#)): $\|\mathbf{x}\|_2 = \sqrt{x^2 + y^2}$ for $\mathbf{x} \in \mathbb{R}^2$



Optimization Nonlinear Optimization

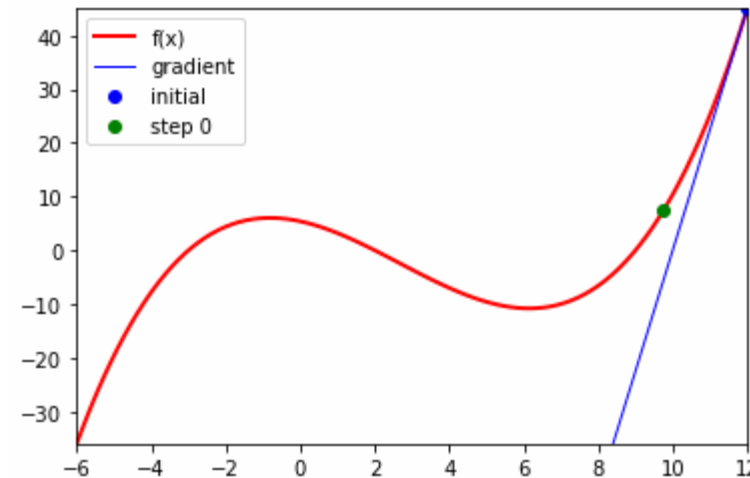
- Nonlinear optimization is the process of solving an optimization problem where some of the constraints or the objective function are **nonlinear**.
 - Alias: Nonlinear programming (NLP)
 - Mathematically, $\hat{\mathbf{x}} = \underset{\mathbf{x}}{\operatorname{argmin}} f(\mathbf{x})$ subject to $g_i(\mathbf{x}) \leq 0$ for each $i \in \{1, \dots, m\}$
 $h_j(\mathbf{x}) = 0$ for each $j \in \{1, \dots, p\}$
 $\mathbf{x} \in X$ (X is a subset of \mathbb{R}^n)
 - $f(\mathbf{x})$: The real-valued objective function
 - $g_i(\mathbf{x})$: The i -th real-valued inequality constraint function
 - $h_j(\mathbf{x})$: The j -th real-valued equality constraint function
 - Example) The objective function $f(x, y) = 4 - (x^2 + y^2)$ is nonlinear.



Nonlinear Optimization

- Gradient descent

- A **first-order iterative algorithm** for finding a local minimum of a differentiable function by pursuing to the opposite direction of the gradient of the function at the current point
- Mathematically, $x_{t+1} = x_t - \gamma f'(x_t)$
 - γ : The step size (a.k.a. learning rate)



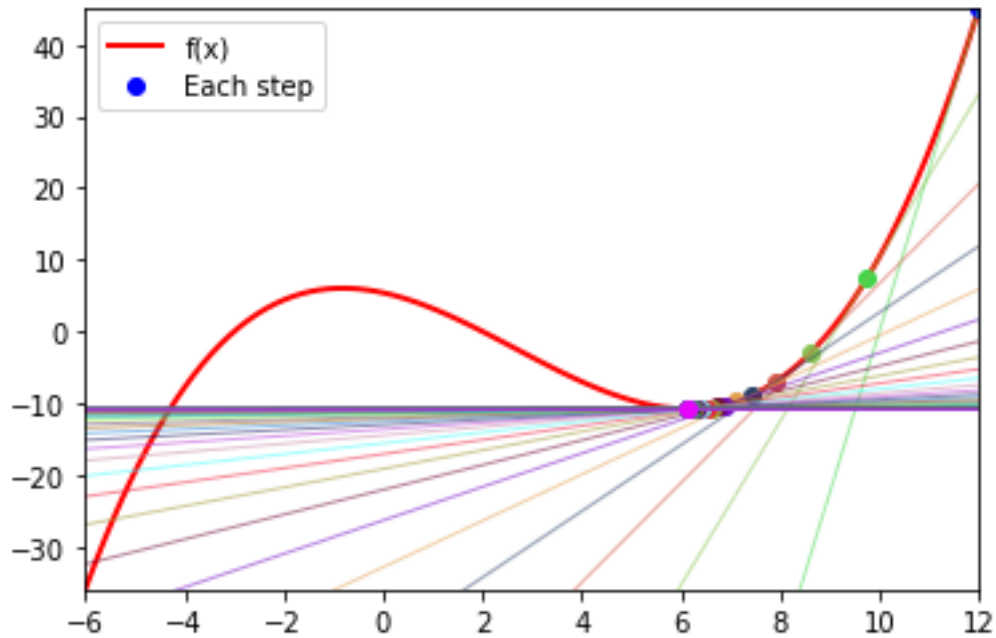
- cf. Stochastic gradient descent (SGD)

- SGD uses an approximated gradient (calculated from a randomly selected subset of the given data) instead of the actual gradient (calculated from the entire data).
- SGD variants: AdaGrad, RMSProp, Adam, ...

Nonlinear Optimization

- Gradient descent

- Example) Find a local minimum $y = 0.1x^3 - 0.8x^2 - 1.5x + 5.4$ from $x = 12$




```

import numpy as np
import matplotlib.pyplot as plt

f = lambda x: 0.1*x**3 - 0.8*x**2 - 1.5*x + 5.4
fd = lambda x: 0.3*x**2 - 1.6*x - 1.5
viz_range = np.array([-6, 12])
learn_rate = 0.1 # Try 0.001, 0.01, 0.5, and 0.6
max_iter = 100
min_tol = 1e-6
x_init = 12 # Try -2

# Prepare visualization
xs = np.linspace(*viz_range, 100)
plt.plot(xs, f(xs), 'r-', label='f(x)', linewidth=2)
plt.plot(x_init, f(x_init), 'b.', label='Each step', markersize=12)
plt.axis((*viz_range, *f(viz_range)))
plt.legend()

x = x_init
for i in range(max_iter):
    # Run the gradient descent
    xp = x
    x = x - learn_rate*fd(x)

    # Update visualization for each iteration
    print(f'Iter: {i}, x = {xp:.3f} to {x:.3f}, f(x) = {f(xp):.3f} to {f(x):.3f} (f\'(x) = {fd(xp):.3f})')
    lcolor = np.random.rand(3)
    approx = fd(xp)*(x-xp) + f(xp)
    plt.plot(xs, approx, '-', linewidth=1, color=lcolor, alpha=0.5)
    plt.plot(x, f(x), '.', color=lcolor, markersize=12)

    # Check the terminal condition
    if abs(x - xp) < min_tol:
        break;


plt.show()

```


Nonlinear Optimization

■ Gradient descent


- Example) Find a local minimum $y = 0.1x^3 - 0.8x^2 - 1.5x + 5.4$ from $x = 12$
 - **Iter: 57**, $x = 6.147$ to 6.147 , $f(x) = -10.822$ to -10.822 ($f'(x) = 0.000$)




Too small step size



Too large step size




Negative infinity



Nonlinear Optimization

▪ Newton's method

- A **second-order iterative algorithm** for finding a local minimum of a differentiable function by pursuing the minima of the locally approximated parabola of the function at the current point
- Mathematically, $x_{t+1} = x_t - \frac{f'(x_t)}{f''(x_t)}$
 - The step size is **not** required.



▪ cf. Gauss-Newton method

- A special case for non-linear least squares problems
 - When the function has a form of $f(x) = r^2(x)$,
 - Newton's method becomes $x_{t+1} = x_t - \frac{r(x_t)}{r'(x_t)}$ (without the 2nd-order derivative)

Nonlinear Optimization

- Newton's method

- Why this equation with the 2nd-order derivative?

- The tangent parabola: The 2nd-order Taylor series expansion at $(x_t, f(x_t))$


$$p(x) = \frac{1}{2}f''(x_t)(x - x_t)^2 + f'(x_t)(x - x_t) + f(x_t)$$

- cf. The tangent line: The 1st-order Taylor series expansion at $(x_t, f(x_t))$

$$l(x) = f'(x_t)(x - x_t) + f(x_t)$$

- Finding the extrema (root) of the tangent parabola, $p'(x) = 0$

$$x_{\text{root}} = x_t - \frac{f'(x_t)}{f''(x_t)}$$

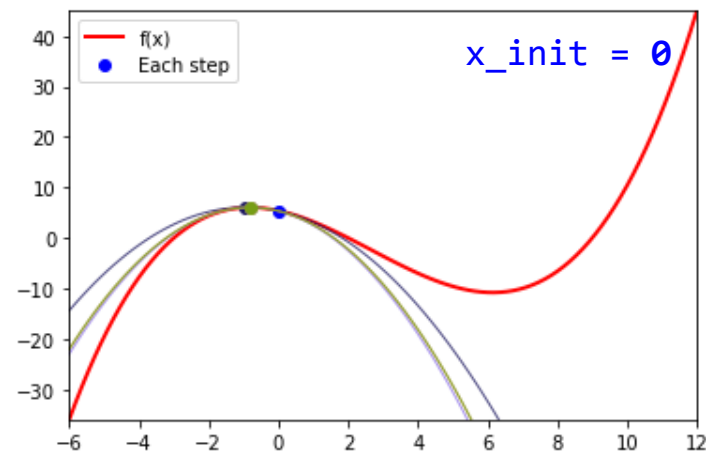
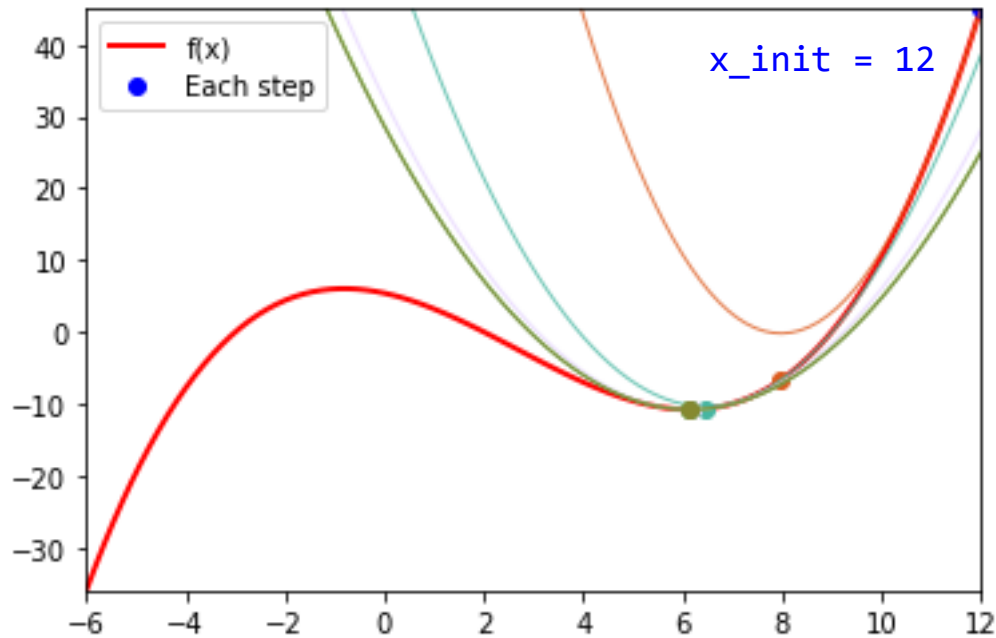


Nonlinear Optimization

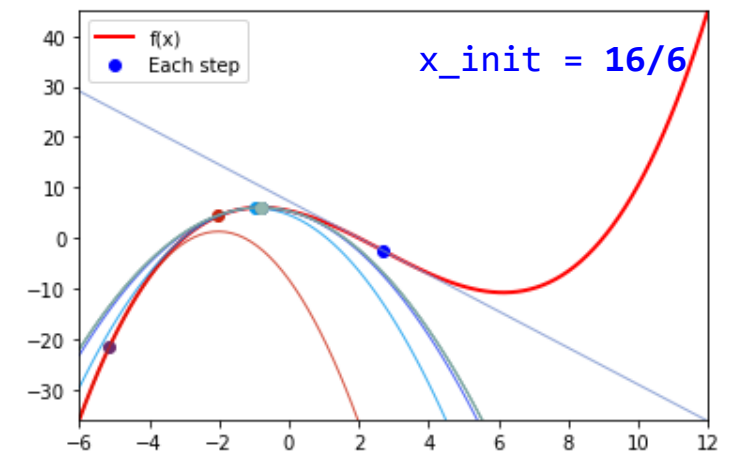
■ Newton's method

– Example) Find a local minimum $y = 0.1x^3 - 0.8x^2 - 1.5x + 5.4$ from $x = 12$

- Iter: 57, $x = 6.147$ to 6.147 , $f(x) = -10.822$ to -10.822 ($f'(x) = 0.000$) # GD
- **Iter: 5**, $x = 6.147$ to 6.147 , $f(x) = -10.822$ to -10.822 (\dots , $f''(x) = 2.088$) # **Newton**



The maxima problem
($f'(x) = 0$ with $f''(x) < 0$)



The saddle point problem
($f''(x) = 0$)

```

import numpy as np
import matplotlib.pyplot as plt

f = lambda x: 0.1*x**3 - 0.8*x**2 - 1.5*x + 5.4
fd = lambda x: 0.3*x**2 - 1.6*x - 1.5
fdd = lambda x: 0.6*x - 1.6
viz_range = np.array([-6, 12])
max_iter = 100
min_tol = 1e-6
x_init = 12 # Try -2, 0, and 16/6 (a saddle point)

# Prepare visualization
xs = np.linspace(*viz_range, 100)
plt.plot(xs, f(xs), 'r-', label='f(x)', linewidth=2)
plt.plot(x_init, f(x_init), 'b.', label='Each step', markersize=12)
plt.axis((*viz_range, *f(viz_range)))
plt.legend()

x = x_init
for i in range(max_iter):
    # Run the Newton method
    xp = x
    x = x - fd(x) / fdd(x) # Replace the denominator as abs(fdd(x)) and (abs(fdd(x)) + 1) to resolve the maxima and saddle point problems

    # Update visualization for each iteration
    print(f'Iter: {i}, x = {xp:.3f} to {x:.3f}, f(x) = {f(xp):.3f} to {f(x):.3f} (f\'(x) = {fd(xp):.3f}, f\'\'(x) = {fdd(xp):.3f})')
    lcolor = np.random.rand(3)
    approx = 0.5*fdd(xp)*(xs-xp)**2 + fd(xp)*(xs-xp) + f(xp)
    plt.plot(xs, approx, '-', linewidth=1, color=lcolor, alpha=0.8)
    plt.plot(x, f(x), '.', color=lcolor, markersize=12)

    # Check the terminal condition
    if abs(x - xp) < min_tol:
        break;

plt.show()

```

scipy.optimize: Optimization and Root Finding

- It provides functions for various optimization problems (and root finding).

- Reference: [Documentation](#) and [Tutorials](#)

- Example) Find a local minimum $y = 0.1x^3 - 0.8x^2 - 1.5x + 5.4$ from $x = 12$ using [scipy.optimize](#)

```
import numpy as np
import matplotlib.pyplot as plt
from scipy.optimize import minimize
```

```
f = lambda x: 0.1*x**3 - 0.8*x**2 - 1.5*x + 5.4
```

```
viz_range = np.array([-6, 12])
```

```
max_iter = 100
```

```
min_tol = 1e-6
```

```
x_init = 12 # Try -2, 0, and 16/6
```

```
# Find the minimum by SciPy
```

```
result = minimize(f, x_init, tol=min_tol, options={'maxiter': max_iter, 'return_all': True})
```

```
print(result)
```

```
# Visualize all iterations
```

```
xs = np.linspace(*viz_range, 100)
```

```
plt.plot(xs, f(xs), 'r-', label='f(x)', linewidth=2)
```

```
xr = np.vstack(result.allvecs)
```

```
plt.plot(xr, f(xr), 'b.', label='Each step', markersize=12)
```

```
plt.legend()
```


```
plt.axis((*viz_range, *f(viz_range)))
```

```
plt.show()
```

- We don't need to provide derivatives.
- We can control its optimization results using parameters (e.g. `tol` and `options`).

scipy.optimize: Optimization and Root Finding

- Example) Find a local minimum $y = 0.1x^3 - 0.8x^2 - 1.5x + 5.4$ from $x = 12$ using [scipy.optimize](#)
 - Iter: 57, $x = 6.147$ to 6.147 , $f(x) = -10.822$ to -10.822 ($f'(x) = 0.000$) # GD
 - Iter: 5, $x = 6.147$ to 6.147 , $f(x) = -10.822$ to -10.822 (\dots , $f''(x) = 2.088$) # Newton
 - allvecs: [array([12.]), array([10.99]), array([8.63764627]), ...] # SciPy
fun: -10.822173403490742
hess_inv: array([[0.47882767]])
jac: array([0.])
message: 'Optimization terminated successfully.'
nfev: 18
nit: 8
njev: 9
status: 0
success: True
x: array([6.14676882])



Objective Functions

- Line representation: $y = wx + b$
- Algebraic distance $d_a = y - (wx_i + b)$
(signed distance)
- Line fitting using $\hat{\mathbf{x}} = \mathbf{A}^+ \mathbf{b} \rightarrow \hat{\mathbf{x}} = \underset{\mathbf{x}}{\operatorname{argmin}} \|\mathbf{A}\mathbf{x} - \mathbf{b}\|_2^2$




(x_i, y_i)

Which line is more closer to the point?

Objective Functions

- Line representation: $ax + by + c = 0$
- Geometric distance $d_g = \frac{ax_i + by_i + c}{\sqrt{a^2 + b^2}} = \frac{\mathbf{n}^T \mathbf{x}_i + c}{\|\mathbf{n}\|}$
(signed distance)



Which line is more closer to the point?

Selecting an objective function (~ a loss function) is important!

Objective Functions

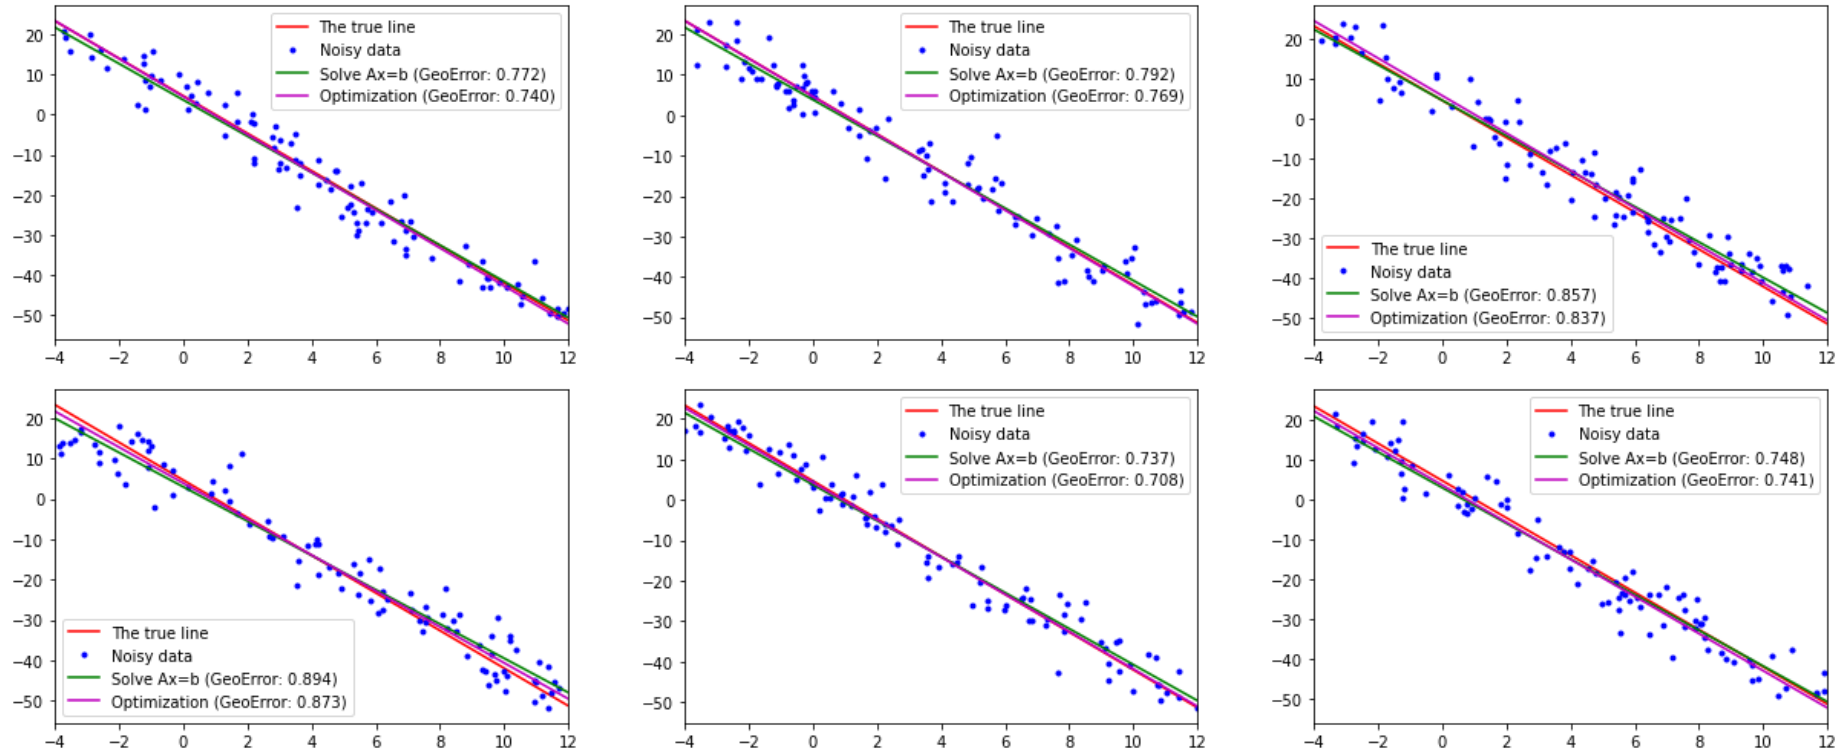
- Example) Line fitting with minimizing **algebraic distance**

– $\hat{\mathbf{x}} = \mathbf{A}^+ \mathbf{b} \rightarrow \hat{\mathbf{x}} = \underset{\mathbf{x}}{\operatorname{argmin}} \|\mathbf{A}\mathbf{x} - \mathbf{b}\|_2^2 = \underset{\mathbf{x}}{\operatorname{argmin}} \sum_i (wx_i + b - y_i)^2$ where $\mathbf{x} = [w, b]$

- Example) Line fitting with minimizing **geometric distance** using [scipy.optimize](#)

– $\hat{\mathbf{x}} = \underset{\mathbf{x}}{\operatorname{argmin}} \sum_i \frac{(wx_i + b - y_i)^2}{w^2 + 1}$ where $\mathbf{x} = [w, b]$ for $ax - y + b = 0$

- cf. Geometric distance will become more helpful when a line has more steeper slope **w**.



```

import numpy as np
import matplotlib.pyplot as plt
from scipy.optimize import minimize

true_line = lambda x: -14/3*x + 14/3
data_range = np.array([-4, 12])
data_num = 100
noise_std = 1

# Generate the true data
x = np.random.uniform(data_range[0], data_range[1], size=data_num)
y = true_line(x)

# Add Gaussian noise
xn = x + np.random.normal(scale=noise_std, size=x.shape)
yn = y + np.random.normal(scale=noise_std, size=y.shape)

# Find a line minimizing algebraic distance
A = np.vstack((xn, np.ones(xn.shape))).T
b = yn
l_alg = np.matmul(np.linalg.pinv(A), b)
e_alg = np.mean(np.abs(l_alg[0]*xn - yn + l_alg[1]) / np.sqrt(l_alg[0]**2 + 1))

# Find a line minimizing geometric distance
geo_dist2 = lambda x: np.sum((x[0]*xn - yn + x[1])**2) / (x[0]**2 + 1)
result = minimize(geo_dist2, [-1, 0]) # The initial value: y = -x
l_geo = result.x
e_geo = np.mean(np.abs(l_geo[0]*xn - yn + l_geo[1]) / np.sqrt(l_geo[0]**2 + 1))

# Plot the data and result
plt.plot(data_range, true_line(data_range), 'r-', label='The true line')
plt.plot(xn, yn, 'b.', label='Noisy data')
plt.plot(data_range, l_alg[0]*data_range + l_alg[1], 'g-', label=f'Solve Ax=b (GeoError: {e_alg:.3f})')
plt.plot(data_range, l_geo[0]*data_range + l_geo[1], 'm-', label=f'Optimization (GeoError: {e_geo:.3f})')
plt.legend()
plt.xlim(data_range)
plt.show()

```

$$\text{cf. } \hat{\mathbf{x}} = \underset{\mathbf{x}}{\operatorname{argmin}} \sum_i (wx_i + b - y_i)^2$$

$$\text{cf. } \hat{\mathbf{x}} = \underset{\mathbf{x}}{\operatorname{argmin}} \sum_i \frac{(wx_i + b - y_i)^2}{w^2 + 1}$$

Summary

■ ~~Optimization~~ Nonlinear optimization

- ~ Finding arguments \mathbf{x} to minimize the *nonlinear* objective function $f(\mathbf{x}) \sim \hat{\mathbf{x}} = \underset{\mathbf{x}}{\operatorname{argmin}} f(\mathbf{x})$
- Gradient descent using the **1st-order** approximation and the given step size.
 - Possible problems: Too small step size, too large step size
 - cf. Stochastic gradient descent (SGD) uses gradient values derived from randomly selected data.
- Newton's method using the **2nd-order** approximation without the step size.
 - Possible problems: The maxima problem, the saddle point problem
 - cf. Gauss-Newton method is a special case for $f(\mathbf{x}) = r^2(\mathbf{x})$.
- scipy.optimize: A sub-module in Scipy for optimization without derivatives
 - You can find minima of any given functions without derivatives.
- **Selecting an objective function is important.**
 - e.g. Algebraic distance vs. geometric distance in line fitting

Our life is full of optimization problems. What is your objective in your life?